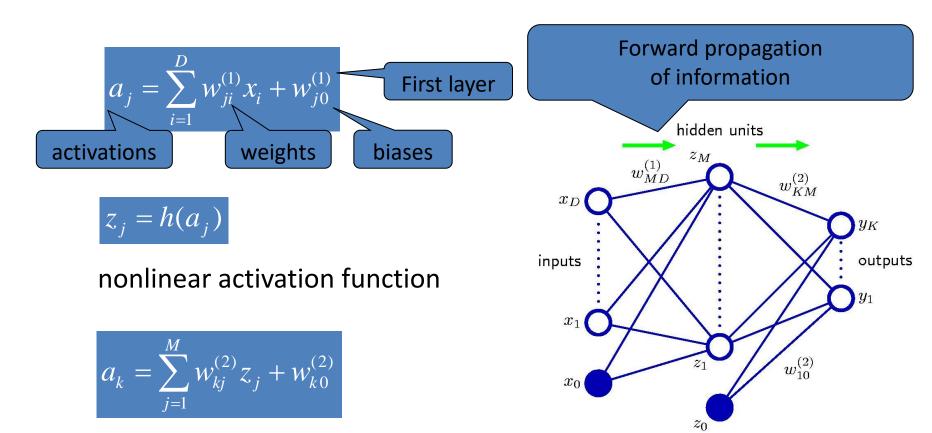


Outlines

- Feedforward Network Functions
- Network Training
- Error Backpropagation
- > Jacobian Matrix
- Hessian Matrix
- CNN and GAN

Feed-forward Network Functions

 Goal: to extend linear model by making the basis functions depend on parameters, allow these parameters to be adjusted.



Outlines

- > Feedforward Network Functions
- Network Training
- > Error Backpropagation
- > Jacobian Matrix
- Hessian Matrix
- CNN and GAN

Network Training

t has a Gaussian distribution with an x-dependent mean

$$p(t \mid x, w) = N(t \mid y(x, w), \beta^{-1})$$

$$p(\mathbf{t} \mid \mathbf{x}, \mathbf{w}, \boldsymbol{\beta}) = \prod_{n=1}^{N} p(t_n \mid x_n, w, \boldsymbol{\beta})$$

(likelihood function)

$$\frac{\beta}{2} \sum_{n=1}^{N} \{y(X_n, w) - t_n\}^2 - \frac{N}{2} \ln \beta + \frac{N}{2} \ln(2\pi)$$
 (negative log)

 Maximizing the likelihood function is equivalent to minimizing the sum-of-squares error function

Network Training

 The choice of output unit activation function and matching error function

Standard regression problems:

Error function: Negative log-likelihood function

Output: identity

For multiple binary classification problems:

Error function: cross-entropy error function

$$E(w) = -\sum_{n=1}^{N} \left\{ t_{n} \ln y_{n} + (1 - t_{n}) \ln (1 - y_{n}) \right\}$$

Output: logistic sigmoid

For multiclass problems:

Multiclass cross-entropy error function a softmax activation function

$$E(w) = -\sum_{n=1}^{N} \sum_{k=1}^{K} t_{nk} \ln y_{nk}$$

$$\frac{\exp(a_k)}{\sum_{i} \exp(a_i)}$$

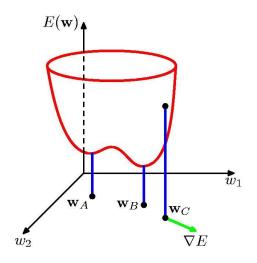
Parameter Optimization

 Error E(w) is a smooth continuous function of w and smallest value will occur at a point in weight space

$$\nabla E(w) = 0$$

Global minimum

Local minima



• Most techniques involve choosing some initial value for weight vector and then moving through weight space in a succession of steps of the form (Many algorithms make use of gradient info)

• Most techniques involve choosing some initial value

• $w^{(\tau+1)} = w^{(\tau)} + \Delta w^{(\tau)}$

Local Quadratic Approximation

Taylor expansion of E(w) around some point

$$E(w) \square E(\hat{w}) + (w - \hat{w})^T \mathbf{b} + \frac{1}{2} (w - \hat{w})^T \mathbf{H} (w - \hat{w})$$

$$\mathbf{b} \equiv \nabla E \big|_{w = \hat{w}} \quad \text{(gradient)} \qquad (\mathbf{H})_{ij} \equiv \frac{\partial E}{\partial w_i \partial w_j} \big|_{w = \hat{w}} \quad \text{(Hessian Matrix)}$$

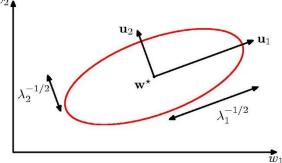
Local approximation to the gradient

$$\nabla E \square \mathbf{b} + \mathbf{H}(w - \hat{w})$$

• Local quadratic approximation when $\nabla E = 0$ at w^{*}

$$E(w) = E(w^*) + \frac{1}{2}(w - w^*)^T \mathbf{H}(w - w^*)$$

$$E(w) = E(w^*) + \frac{1}{2} \sum_{i} \lambda_i \alpha_i^2$$



Use of Gradient Information

- In the quadratic approximation, computational cost to find minimum is O(W³)
 - W is the dimensionality of w
 - perform O(W²) evaluations, each of which would require O(W) steps.
- In an algorithm that makes use of the gradient information, computational cost is O(W²)
 - By using error backpropagation, O(W) gradient evaluations and each such evaluation takes only O(W) steps.

Gradient Descent Optimization

 Weight update to comprise a small step in the direction of the negative gradient

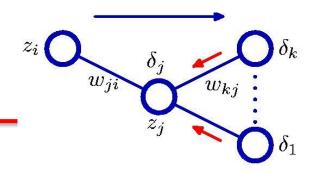
$$w^{(\tau+1)} = w^{(\tau)} - \eta \nabla E(w^{(\tau)})$$

- Batch method
 - Techniques that use the whole data set at once.
 - The error function always decreases at each iteration unless the weight vector has arrived at a local or global minimum.
- On-line version of gradient descent
 - Sequential gradient descent or stochastic gradient descent
 - Update to the weight vector based on one data point at a time.
 - Can handle redundancy in the data much more efficiently.
 - The possibility of escaping from local minima.

Outlines

- > Feedforward Network Functions
- Network Training
- Error Backpropagation
- > Jacobian Matrix
- Hessian Matrix
- CNN and GAN

Error Backpropagation



Error Backpropagation

Apply an input vector X_n to the network and forward propagate through the network using equations below to find the activations of all the hidden and output units.

$$a_j = \sum_i w_{ji} z_i$$

$$z_j = h(a_j)$$

Evaluate the δ_k for all the output units using

$$\delta_k = y_k - t_k$$

Backpropagate the δ 's using

for each hidden unit in the network

$$\delta_j = h'(a_j) \sum_k w_{ji} \delta_k$$

Use $\frac{\partial Z_n}{\partial w} = \delta_j Z_i$ to evaluate the required derivatives.

A Simple Example

$$h(a) \equiv \tanh(a)$$

$$h'(a) = 1 - h(a)^2$$

$$E_n = \frac{1}{2} \sum_{k=1}^{K} (y_k - t_k)^2$$

$$a_{j} = \sum_{i=1}^{D} w_{ji}^{(1)} x_{i}$$

$$z_j = \tanh(a_j)$$

$$y_k = \sum_{j=0}^{M} w_{kj}^{(2)} z_j$$

$$tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

(y_k: output unit k, t_k: the corresponding target)

$$\delta_j = (1 - z_j^2) \sum_{k=1}^K w_{kj} \delta_k$$

$$\frac{\partial E_n}{\partial w_{ji}^{(1)}} = \mathcal{S}_j x_i$$

$$\frac{\partial E_n}{\partial w_{kj}^{(2)}} = \delta_k z_j$$

Outlines

- > Feedforward Network Functions
- Network Training
- > Error Backpropagation
- Jacobian Matrix
- Hessian Matrix
- CNN and GAN

The Jacobian Matrix

- The technique of backpropagation can also be applied to the calculation of other derivatives.
- The Jacobian matrix
 - Elements are given by the derivatives of the network outputs w.r.t. the input:

$$J_{ki} = \frac{\partial y_k}{\partial x_i}$$

Minimizing an error function E w.r.t. the parameter

$$\frac{\partial E}{\partial w} = \sum_{k,j} \frac{\partial E}{\partial y_k} \frac{\partial y_k}{\partial z_j} \frac{\partial z_j}{\partial w}$$

The Jacobian Matrix

- A measure of the *local sensitivity* of the outputs to change in each of the input variables.
 - In general, the network mapping is nonlinear, and so the elements will not be constants. This is valid provided Δx_i are small.

$$\Delta y_k \, \Box \, \sum_i \frac{\partial y_k}{\partial x_j} \, \Delta x_i$$

The evaluation of the Jacobian Matrix

$$\frac{\partial y_k}{\partial a_l} = \delta_{kj} \sigma'(a_j)$$

$$J_{ki} = \frac{\partial y_k}{\partial x_i} = \sum_j \frac{\partial y_k}{\partial a_j} \frac{\partial a_j}{\partial x_i} = \sum_j w_{ji} \frac{\partial y_k}{\partial a_j} \qquad \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_j} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_j} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_j} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_j} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial x_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial x_l}{\partial a_j} = h'(a_j) \sum_l w_{lj} \frac{\partial y_k}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial x_l}{\partial a_j} = \sum_l \frac{\partial x_l}{\partial a_l} \frac{\partial x_l}{\partial a_j} = \sum_l \frac{\partial y_k}{\partial a_l} \frac{\partial x_l}{\partial a_j} = \sum_l \frac{\partial x_l}{\partial a_l} \frac{\partial x_l}{\partial a_j} = \sum_l \frac{\partial x_l}{\partial a_l} \frac{\partial x_l}{\partial a_l} = \sum_l$$

$$\frac{\partial y_k}{\partial a_j} = \sum_{l} \frac{\partial y_k}{\partial a_l} \frac{\partial a_l}{\partial a_j} = h'(a_j) \sum_{l} w_{lj} \frac{\partial y_k}{\partial a_l}$$

(Sigmoidal activation function)

Outlines

- > Feedforward Network Functions
- Network Training
- Error Backpropagation
- > Jacobian Matrix
- Hessian Matrix
- CNN and GAN

Hessian: Diagonal approximation

- Inverse of the Hessian is useful, and inverse of diagonal matrix is trivial to evaluate.
- Consider an error function
 - Replaces the off-diagonal elements with zeros
 - The diagonal elements of the Hessian:

$$\frac{\partial^{2} E_{n}}{\partial w_{ji}^{2}} = \frac{\partial^{2} E_{n}}{\partial a_{j}^{2}} z_{i}^{2}$$

$$\frac{\partial^{2} E_{n}}{\partial a_{j}^{2}} = h'(a_{j})^{2} \sum_{k} \sum_{k'} w_{kj} w_{k'j} \frac{\partial^{2} E_{n}}{\partial a_{k} \partial a_{k'}} + h''(a_{j}) \sum_{k} w_{kj} \frac{\partial E_{n}}{\partial a_{k}} \frac{\partial E_{n}}{\partial a_{k}}$$

$$\frac{\partial^2 E_n}{\partial a_j^2} = h'(a_j)^2 \sum_k w_{kj}^2 \frac{\partial^2 E_n}{\partial a_k^2} + h''(a_j) \sum_k w_{kj} \frac{\partial E_n}{\partial a_k}$$

(neglect off-diagonal)

Outer Product Approximation

$$\mathbf{H} = \nabla \nabla E = \sum_{n=1}^{N} \nabla y_n \nabla y_n + \sum_{n=1}^{N} (y_n - t_n) \nabla \nabla y_n$$

- Output y happen to be very close to the target values t, the second term will be small and can be neglected.
 - Or the value of y t is uncorrelated with the value of the second derivative term, then the whole term will average to zero.

$$\mathbf{H} \square \sum_{n=1}^{N} \mathbf{b}_{n} \mathbf{b}_{n}^{T}$$

$$\mathbf{H} \square \sum_{n} \mathbf{b}_{n} \mathbf{b}_{n}^{T} \qquad \mathbf{b}_{n} = \nabla y_{n} = \nabla a_{n}$$

Inverse Hessian

- A procedure for approximating the inverse of the Hessian
 - First, we write the outer-product approximation
 - Derive a sequential procedure for building up the Hessian by including data points one at a time.

$$\mathbf{H}_{L+1} = \mathbf{H}_{L} + \mathbf{b}_{L+1} \mathbf{b}_{L+1}^{T}$$
 $\mathbf{H}_{L+1}^{-1} = \mathbf{H}_{L}^{-1} - \frac{\mathbf{H}_{L}^{-1} \mathbf{b}_{L+1} \mathbf{b}_{L+1}^{T} \mathbf{H}_{L}^{-1}}{1 + \mathbf{b}_{L+1}^{T} \mathbf{H}_{L}^{-1} \mathbf{b}_{L+1}}$

• The initial matrix $H_0 = H + \alpha I$ α is a small quantity, not sensitive to the precise value of α .

Finite Differences

 By using a symmetrical central differences formulation

$$\frac{\partial^{2} E}{\partial w_{ji} \partial w_{lk}} = \frac{1}{4\varepsilon} \{ E(w_{ji} + \varepsilon, w_{lk} + \varepsilon) - E(w_{ji} + \varepsilon, w_{lk} - \varepsilon)$$
$$-E(w_{ji} - \varepsilon, w_{lk} + \varepsilon) + E(w_{ji} - \varepsilon, w_{lk} - \varepsilon) \} + O(\varepsilon^{2})$$

Require O(W³) operations to evaluate the complete Hessian

$$\frac{\partial^2 E}{\partial w_{ji} \partial w_{lk}} = \frac{1}{2\varepsilon} \left\{ \frac{\partial E}{\partial w_{ji}} (w_{lk} + \varepsilon) - \frac{\partial E}{\partial w_{ji}} (w_{lk} - \varepsilon) \right\} + O(\varepsilon^2)$$

By applying central differences to the first derivatives of the error function. Costs: O(W²)

Evaluation of the Hessian

The Hessian can also be evaluated exactly.

Using extension of the technique of backpropagation used to evaluate first derivatives.

$$M_{kk'} = rac{\partial^2 E_n}{\partial a_k \partial a_{k'}}$$

(Both weights in the second layer)

 $\frac{\partial^2 E_n}{\partial w_{kj}^{(2)} \partial w_{k'j'}^{(2)}} = z_j z_j M_{kk'} \qquad \delta_k = \frac{\partial E_n}{\partial a_k}$

$$\frac{\partial^2 E_n}{\partial w_{ji}^{(1)} \partial w_{j'i'}^{(1)}} = x_i x_{i'} h''(a_{j'}) I_{jj'} \sum_k w_{kj'}^{(2)} \delta_k$$

$$+x_i x_i h'(a_{j'})h'(a_j) \sum_{k} \sum_{k'} w_{k'j'}^{(2)} w_{kj}^{(2)} M_{kk'}$$

$$\frac{\partial^2 E_n}{\partial w_{ji}^{(1)} \partial w_{kj'}^{(2)}} = x_i h'(a_{j'}) \{ \delta_k I_{jj'} + z_j \sum_{k'} w_{k'j'}^{(2)} H_{kk'} \}$$

(Both weights in each layer)

Fast Multiplication by the Hessian

- Try to find an efficient approach to evaluating v^TH directly. O(W) operations
 - $v^TH = v^T\nabla (\nabla E)$ (∇ : the gradient operator in weight space)
 - Introducing new notation $R\{w\} = v$, to denote the operator $v^T \nabla$

$$R\{a_j\} = \sum_{i} v_{ji} x_i$$
 $R\{z_j\} = h'(a_j)R\{a_j\}$

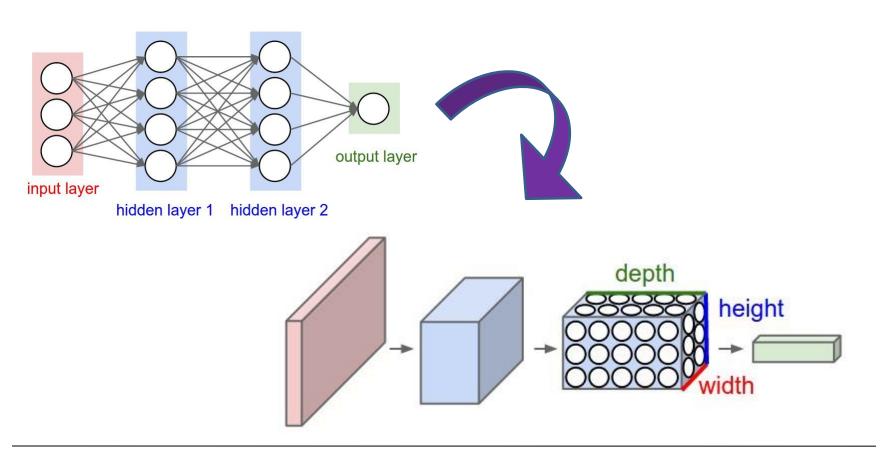
$$R\{\delta_{j}\} = h''(a_{j})R\{a_{j}\} \sum_{k} w_{kj}\delta_{k} + h'(a_{j}) \sum_{k} v_{kj}\delta_{k} + h'(a_{j}) \sum_{k} w_{kj}R\{\delta_{k}\}$$

- The Implementation of this algorithm involves the introduction of additional variables $R\{a_j\}$ $R\{z_j\}$ for the hidden units and $R\{\delta_k\}$ $R\{y_k\}$ for the output units.
- For each input pattern, the values of these quantities can be found.

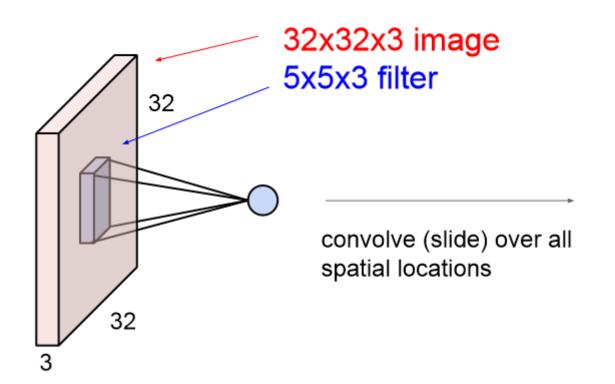
Outlines

- Feedforward Network Functions
- Network Training
- Error Backpropagation
- > Jacobian Matrix
- Hessian Matrix
- CNN and GAN

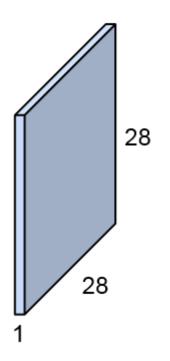
What's the same and difference between a normal full-connected network and convolution network



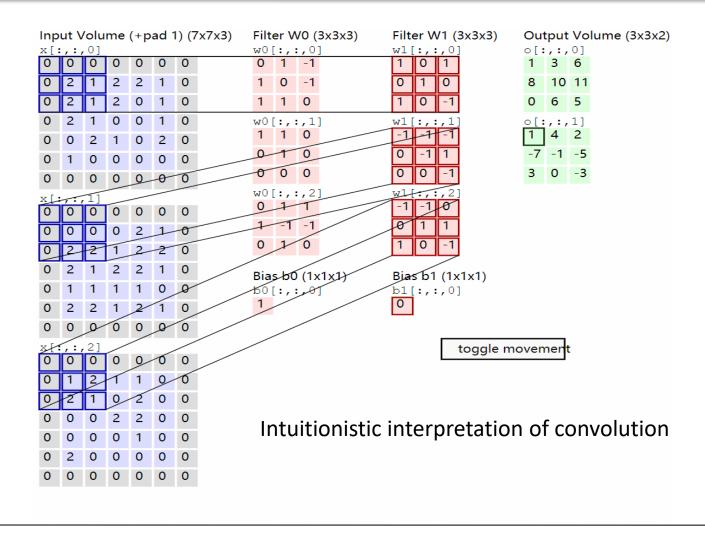
Convolution Layer



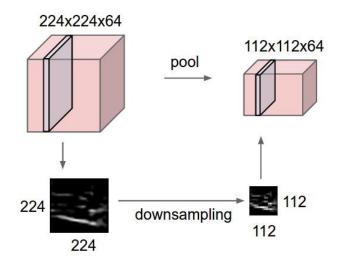
activation map

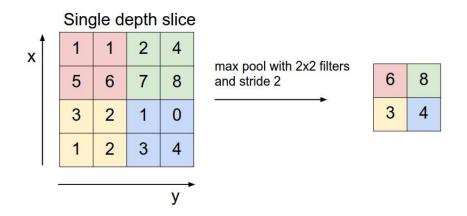


- Convolution layer
- Pooling layer
- Full-connected layer
- Rectified Linear Units (ReLu) layer



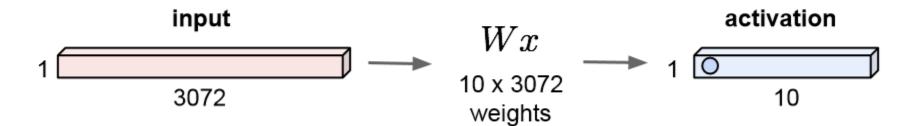
Pooling layer: to reduce the size of the input Usually using the max value in the block(2*2 most usually) to replace the block itself

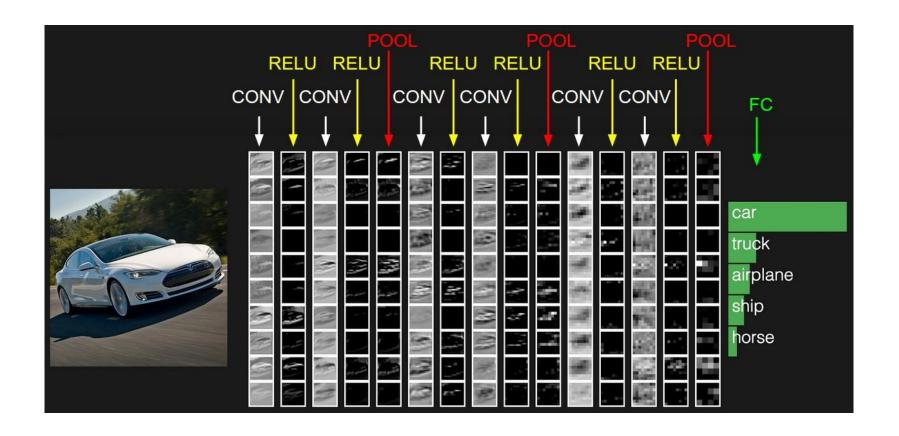




Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1





CNN

Useful page

CS231n: Convolutional Neural Networks for Visual Recognition:

http://cs231n.github.io/

Neural Networks and Deep Learning

http://neuralnetworksanddeeplearning.com/

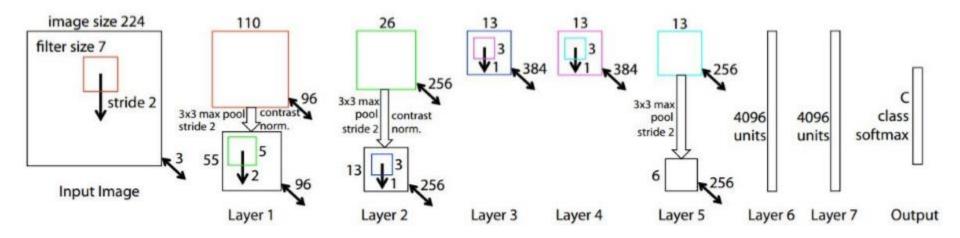
CVPR 2017 Paper interpretation

http://cvmart.net/community/article/detail/69

VGG

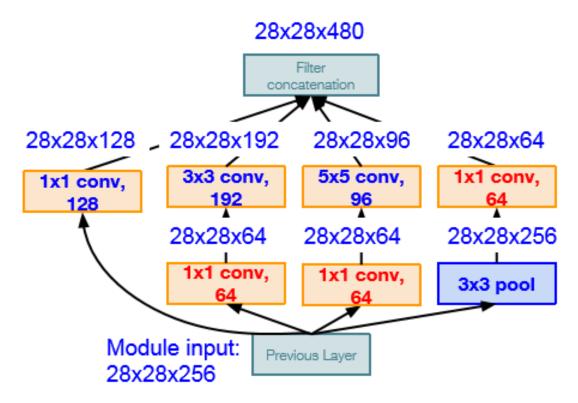
```
(not counting biases)
INPUT: [224x224x3]
                      memory: 224*224*3=150K params: 0
CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*3)*64 = 1.728
                                                                                               FC 1000
                                                                                                          fc8
CONV3-64: [224x224x64] memory: 224*224*64=3.2M params: (3*3*64)*64 = 36,864
                                                                                               FC 4098
                                                                                                          fc7
POOL2: [112x112x64] memory: 112*112*64=800K params: 0
                                                                                               FC 4098
                                                                                                          fc6
CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*64)*128 = 73,728
CONV3-128: [112x112x128] memory: 112*112*128=1.6M params: (3*3*128)*128 = 147,456
                                                                                                        conv5-3
POOL2: [56x56x128] memory: 56*56*128=400K params: 0
                                                                                                        conv5-2
CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*128)*256 = 294,912
                                                                                                        conv5-1
                                                                                                Pool
CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824
                                                                                                        conv4-3
CONV3-256: [56x56x256] memory: 56*56*256=800K params: (3*3*256)*256 = 589,824
                                                                                                        conv4-2
POOL2: [28x28x256] memory: 28*28*256=200K params: 0
                                                                                                        conv4-1
CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*256)*512 = 1,179,648
                                                                                                Pool
CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512=2,359,296
                                                                                                        conv3-2
CONV3-512: [28x28x512] memory: 28*28*512=400K params: (3*3*512)*512 = 2.359,296
                                                                                                        conv3-1
POOL2: [14x14x512] memory: 14*14*512=100K params: 0
                                                                                              Pool
CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296
                                                                                                        conv2-2
                                                                                                        conv2-1
CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2,359,296
CONV3-512: [14x14x512] memory: 14*14*512=100K params: (3*3*512)*512 = 2.359,296
                                                                                                        conv1-2
POOL2: [7x7x512] memory: 7*7*512=25K params: 0
                                                                                                        conv1-1
FC: [1x1x4096] memory: 4096 params: 7*7*512*4096 = 102,760,448
                                                                                                Input
FC: [1x1x4096] memory: 4096 params: 4096*4096 = 16,777,216
                                                                                             VGG16
FC: [1x1x1000] memory: 1000 params: 4096*1000 = 4,096,000
TOTAL memory: 24M * 4 bytes ~= 96MB / image (only forward! ~*2 for bwd)
                                                                                            Common names
TOTAL params: 138M parameters
```

AlexNet/ZFNet



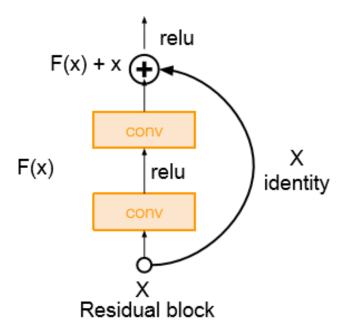
a large, deep convolutional neural network to classify the 1.3 million high-resolution images in the LSVRC-2010 ImageNet training set into the 1000 different classes http://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks

GoogleNet

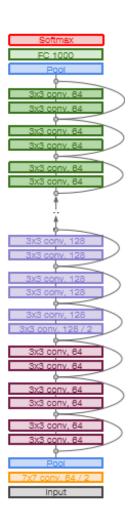


Inception Module reduced a huge number of parameters in the network (4M, compared to AlexNet with 60M). using Average Pooling instead of Fully Connected layers at the top of the ConvNet, eliminating a large amount of parameters that do not seem to matter much. https://arxiv.org/pdf/1409.4842.pdf

ResNet

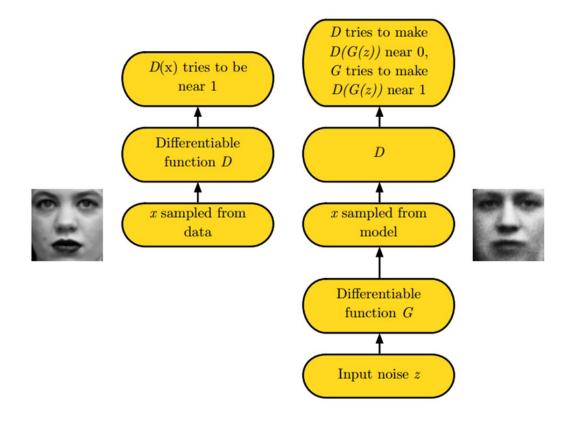


Deep Residual Learning for Image Recognition https://arxiv.org/pdf/1512.03385.pdf



GAN

Generative Adversarial Networks



GAN

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k, is a hyperparameter. We used k = 1, the least expensive option, in our experiments.

for number of training iterations do

for k steps do

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D\left(\boldsymbol{x}^{(i)}\right) + \log\left(1 - D\left(G\left(\boldsymbol{z}^{(i)}\right)\right)\right) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_q(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^{m} \log \left(1 - D \left(G \left(\boldsymbol{z}^{(i)} \right) \right) \right).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

Generative Adversarial Nets

https://arxiv.org/pdf/1406.2661.pdf

HW5

NN Fundamentals: 5.1 5.2 5.4 5.9

Hessian: 5.16 5.19

Extensions: 5.34 5.37 5.38 5.39